

# Neural Networks Deep Learning – Icp8

Name: Deekshitha, Gaddameedhi

ID: 700755765

GitHub link: [https://github.com/deekshitha430/icp8\\_neural](https://github.com/deekshitha430/icp8_neural)

Video Link: [https://drive.google.com/file/d/19v3Dc0a5H-b0bkvmWFZ3v1yG\\_xtfgo7/view?usp=sharing](https://drive.google.com/file/d/19v3Dc0a5H-b0bkvmWFZ3v1yG_xtfgo7/view?usp=sharing)

In class programming:

## 1. Add one more hidden layer to autoencoder

```
In [10]: from keras.layers import Input, Dense
         from keras.models import Model

         # this is the size of our encoded representations
         encoding_dim = 32 # 32 floats -> compression of factor 24.5, assuming the input is 784 floats

         # this is our input placeholder
         input_img = Input(shape=(784,))
         # "encoded" is the encoded representation of the input
         encoded = Dense(encoding_dim, activation='relu')(input_img)
         # "decoded" is the lossy reconstruction of the input
         decoded = Dense(784, activation='sigmoid')(encoded)
         # this model maps an input to its reconstruction
         autoencoder = Model(input_img, decoded)
         # this model maps an input to its encoded representation
         autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy', metrics=['accuracy'])
         from keras.datasets import mnist, fashion_mnist
         import numpy as np
         (x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
         x_train = x_train.astype('float32') / 255.
         x_test = x_test.astype('float32') / 255.
         x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
         x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

         autoencoder.fit(x_train, x_train,
                         epochs=5,
                         batch_size=256,
                         shuffle=True,
                         validation_data=(x_test, x_test))
```

```
shuffle=True,
validation_data=(x_test, x_test))
```

```
Epoch 1/5
235/235 [=====] - 5s 17ms/step - loss: 0.6953 - accuracy: 9.0000e-04 - val_loss: 0.6952 - val_accu
acy: 8.0000e-04
Epoch 2/5
235/235 [=====] - 3s 14ms/step - loss: 0.6951 - accuracy: 9.0000e-04 - val_loss: 0.6949 - val_accu
acy: 8.0000e-04
Epoch 3/5
235/235 [=====] - 3s 14ms/step - loss: 0.6948 - accuracy: 9.1667e-04 - val_loss: 0.6947 - val_accu
acy: 8.0000e-04
Epoch 4/5
235/235 [=====] - 3s 13ms/step - loss: 0.6946 - accuracy: 9.3333e-04 - val_loss: 0.6945 - val_accu
acy: 9.0000e-04
Epoch 5/5
235/235 [=====] - 3s 13ms/step - loss: 0.6944 - accuracy: 9.5000e-04 - val_loss: 0.6943 - val_accu
acy: 9.0000e-04
```

```
: [10]: <keras.src.callbacks.History at 0x23310f8b310>
```

## Neural Networks Deep Learning – Icp8

```
In [11]: from keras.layers import Input, Dense
         from keras.models import Model

         # This is the size of our encoded representation
         encoding_dim = 32 # 32 floats -> compression of factor 24.5, assuming the input is 784 floats

         # This is our input placeholder
         input_img = Input(shape=(784,))

         # "encoded" is the encoded representation of the input
         encoded1 = Dense(128, activation='relu')(input_img)
         encoded2 = Dense(encoding_dim, activation='relu')(encoded1)

         # "decoded" is the lossy reconstruction of the input
         decoded1 = Dense(128, activation='relu')(encoded2)
         decoded2 = Dense(784, activation='sigmoid')(decoded1)

         # This model maps an input to its reconstruction
         autoencoder = Model(input_img, decoded2)

         # This model maps an input to its encoded representation
         encoder = Model(input_img, encoded2)

         # This is our decoder model
         encoded_input = Input(shape=(encoding_dim,))
         decoder_layer1 = autoencoder.layers[-2]
         decoder_layer2 = autoencoder.layers[-1]
         decoder = Model(encoded_input, decoder_layer2(decoder_layer1(encoded_input)))

         # Compile the model
         autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy', metrics=['accuracy'])

         # Load the MNIST dataset
         from keras.datasets import mnist, fashion_mnist
         import numpy as np
         (x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()

         # Normalize and flatten the data
         x_train = x_train.astype('float32') / 255.
         x_test = x_test.astype('float32') / 255.
         x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
         x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

         # Train the autoencoder
         autoencoder.fit(x_train, x_train,
                         epochs=5,
                         batch_size=256,
                         shuffle=True,
                         validation_data=(x_test, x_test))

         encoder.save('encoder.h5')
         decoder.save('decoder.h5')

Epoch 1/5
235/235 [=====] - 7s 21ms/step - loss: 0.6944 - accuracy: 0.0011 - val_loss: 0.6943 - val_accuracy: 0.0016
Epoch 2/5
235/235 [=====] - 4s 17ms/step - loss: 0.6943 - accuracy: 0.0011 - val_loss: 0.6942 - val_accuracy: 0.0016
Epoch 3/5
235/235 [=====] - 4s 17ms/step - loss: 0.6941 - accuracy: 0.0011 - val_loss: 0.6940 - val_accuracy: 0.0016
Epoch 4/5
235/235 [=====] - 4s 17ms/step - loss: 0.6940 - accuracy: 0.0011 - val_loss: 0.6939 - val_accuracy: 0.0015
Epoch 5/5
235/235 [=====] - 4s 17ms/step - loss: 0.6939 - accuracy: 0.0011 - val_loss: 0.6938 - val_accuracy: 0.0015

it[11]: <keras.src.callbacks.History at 0x23310fb19d0>
```

## Neural Networks Deep Learning – Icp8

2. Do the prediction on the test data and then visualize one of the reconstructed version of that test data. Also, visualize the same test data before reconstruction using Matplotlib

```
[16]: ► import matplotlib.pyplot as plt

# Get the reconstructed images for the test set
reconstructed_imgs = autoencoder.predict(x_test)

# Choose a random image from the test set
n = 10 # index of the image to be plotted
plt.figure(figsize=(10, 5))

# Plot the original image
ax = plt.subplot(1, 2, 1)
plt.imshow(x_test[n].reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
ax.set_title("Original Image")


# Plot the reconstructed image
ax = plt.subplot(1, 2, 2)
plt.imshow(reconstructed_imgs[n].reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
ax.set_title("Reconstructed Image")

plt.show()
```


313/313 [=====] - 1s 3ms/step

313/313 [=====] - 1s 3ms/step

Original Image



Reconstructed Image



## Neural Networks Deep Learning – lcp8

### 3. Repeat the question 2 on the denoising autoencoder

```
n [17]: from keras.layers import Input, Dense
        from keras.models import Model

        # this is the size of our encoded representations
        encoding_dim = 32 # 32 floats -> compression of factor 24.5, assuming the input is 784 floats

        # this is our input placeholder
        input_img = Input(shape=(784,))
        # "encoded" is the encoded representation of the input
        encoded = Dense(encoding_dim, activation='relu')(input_img)
        # "decoded" is the lossy reconstruction of the input
        decoded = Dense(784, activation='sigmoid')(encoded)
        # this model maps an input to its reconstruction
        autoencoder = Model(input_img, decoded)
        # this model maps an input to its encoded representation
        autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy', metrics=['accuracy'])
        from keras.datasets import fashion_mnist
        import numpy as np
        (x_train, _), (x_test, _) = fashion_mnist.load_data()
        x_train = x_train.astype('float32') / 255.
        x_test = x_test.astype('float32') / 255.
        x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
        x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

        #introducing noise
        noise_factor = 0.5
        x_train_noisy = x_train + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_train.shape)
        x_test_noisy = x_test + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_test.shape)

        autoencoder.fit(x_train_noisy, x_train,
                        epochs=10,
                        batch_size=256,
                        shuffle=True,
                        validation_data=(x_test_noisy, x_test_noisy))
```

```
Epoch 1/10
235/235 [=====] - 5s 16ms/step - loss: 0.6982 - accuracy: 5.3333e-04 - val_loss: 0.6979 - val_accu
acy: 8.0000e-04
Epoch 2/10
235/235 [=====] - 3s 13ms/step - loss: 0.6978 - accuracy: 5.3333e-04 - val_loss: 0.6976 - val_accu
acy: 8.0000e-04
Epoch 3/10
235/235 [=====] - 3s 11ms/step - loss: 0.6975 - accuracy: 5.3333e-04 - val_loss: 0.6972 - val_accu
acy: 8.0000e-04
Epoch 4/10
235/235 [=====] - 3s 14ms/step - loss: 0.6971 - accuracy: 5.3333e-04 - val_loss: 0.6969 - val_accu
acy: 9.0000e-04
Epoch 5/10
235/235 [=====] - 3s 13ms/step - loss: 0.6968 - accuracy: 5.5000e-04 - val_loss: 0.6966 - val_accu
acy: 9.0000e-04
Epoch 6/10
235/235 [=====] - 3s 13ms/step - loss: 0.6965 - accuracy: 5.6667e-04 - val_loss: 0.6963 - val_accu
acy: 9.0000e-04
Epoch 7/10
235/235 [=====] - 3s 13ms/step - loss: 0.6962 - accuracy: 6.0000e-04 - val_loss: 0.6960 - val_accu
acy: 9.0000e-04
Epoch 8/10
235/235 [=====] - 3s 13ms/step - loss: 0.6960 - accuracy: 6.1667e-04 - val_loss: 0.6957 - val_accu
acy: 9.0000e-04
Epoch 9/10
235/235 [=====] - 4s 15ms/step - loss: 0.6957 - accuracy: 6.0000e-04 - val_loss: 0.6955 - val_accu
acy: 9.0000e-04
Epoch 10/10
235/235 [=====] - 3s 15ms/step - loss: 0.6954 - accuracy: 6.0000e-04 - val_loss: 0.6952 - val_accu
acy: 9.0000e-04
```

## Neural Networks Deep Learning – lcp8

```
In [20]: import matplotlib.pyplot as plt

# Get the reconstructed images for the test set
reconstructed_imgs = autoencoder.predict(x_test_noisy)

# Choose a random image from the test set
n = 10 # index of the image to be plotted
plt.figure(figsize=(10, 5))


# Plot the original noisy image
ax = plt.subplot(1, 2, 1)
plt.imshow(x_test_noisy[n].reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
ax.set_title("Noisy Image")

# Plot the reconstructed image
ax = plt.subplot(1, 2, 2)
plt.imshow(reconstructed_imgs[n].reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
ax.set_title("Reconstructed Image")

plt.show()

313/313 [=====] - 1s 3ms/step
```

313/313 [=====] - 1s 3ms/step



### 4. plot loss and accuracy using the history object

## Neural Networks Deep Learning – lcp8

```
In [25]: import matplotlib.pyplot as plt

# Train the autoencoder
history = autoencoder.fit(x_train_noisy, x_train,
                          epochs=10,
                          batch_size=256,
                          shuffle=True,
                          validation_data=(x_test_noisy, x_test_noisy))

# Plot the loss
plt.plot(history.history['loss'], label='train')
plt.plot(history.history['val_loss'], label='test')
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend()
plt.show()

# Plot the accuracy
plt.plot(history.history['accuracy'], label='train')
plt.plot(history.history['val_accuracy'], label='test')
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend()
plt.show()
```

```
plt.show()
```

```
Epoch 1/10
235/235 [=====] - 3s 14ms/step - loss: 0.6944 - accuracy: 0.0028 - val_loss: 0.6942 - val_accuracy: 0.0018
Epoch 2/10
235/235 [=====] - 3s 14ms/step - loss: 0.6942 - accuracy: 0.0027 - val_loss: 0.6940 - val_accuracy: 0.0018
Epoch 3/10
235/235 [=====] - 3s 13ms/step - loss: 0.6940 - accuracy: 0.0027 - val_loss: 0.6938 - val_accuracy: 0.0018
Epoch 4/10
235/235 [=====] - 3s 13ms/step - loss: 0.6938 - accuracy: 0.0027 - val_loss: 0.6936 - val_accuracy: 0.0018
Epoch 5/10
235/235 [=====] - 3s 13ms/step - loss: 0.6936 - accuracy: 0.0027 - val_loss: 0.6934 - val_accuracy: 0.0019
Epoch 6/10
235/235 [=====] - 3s 14ms/step - loss: 0.6935 - accuracy: 0.0027 - val_loss: 0.6933 - val_accuracy: 0.0019
Epoch 7/10
235/235 [=====] - 3s 13ms/step - loss: 0.6933 - accuracy: 0.0027 - val_loss: 0.6931 - val_accuracy: 0.0020
Epoch 8/10
235/235 [=====] - 3s 13ms/step - loss: 0.6931 - accuracy: 0.0027 - val_loss: 0.6929 - val_accuracy: 0.0021
Epoch 9/10
235/235 [=====] - 3s 14ms/step - loss: 0.6930 - accuracy: 0.0028 - val_loss: 0.6928 - val_accuracy: 0.0021
Epoch 10/10
235/235 [=====] - 3s 13ms/step - loss: 0.6928 - accuracy: 0.0028 - val_loss: 0.6926 - val_accuracy: 0.0021
```

## Neural Networks Deep Learning – lcp8

